Automated Weed Detection Using CNN-ViT Hybrid for Precision Agriculture

Reducing Herbicide Usage Through Edge-AI Weed Classification

Course : Embedded AI

Mentor : Dr. Dheeraj

Team Members: Deepesh, Jewel, Vaibhav

The Problem

$30B

Global Market

Annual herbicide market value

70%

Wasted Spray

Herbicide hitting soil instead of weeds

40-80%

Potential Reduction

Herbicide savings with precision agriculture

Traditional Farming Challenges

Traditional farming practices spray entire fields uniformly, resulting in massive waste and environmental damage. This approach causes water contamination, increased costs for farmers, and significant pollution.

Precision agriculture goal: Spray only where weeds exist to reduce herbicide usage, lower costs, increase yield, and reduce environmental impact.

Key question: Can we automatically detect weeds in the field using lightweight AI that runs on edge devices?

Project Objective & Dataset

Primary Objective

Automate weed detection using computationally inexpensive CNNs to reduce herbicide usage and increase precision agriculture yield

Classification Goals

- Classify 8 weed species plus negative class
- Binary decision (Weed or Not-Weed)
- Spray decision (No/Spot/Full spray)

Performance Target

Run on NVIDIA Jetson Nano edge device with real-time inference over 15 FPS

DeepWeeds Dataset

Dataset Overview

17,509 color images at 256×256 resolution from Olsen et al., 2019

Collection sites: 8 rangeland locations across Queensland, Australia

Key challenge: Severe 9:1 class imbalance between negative and weed classes

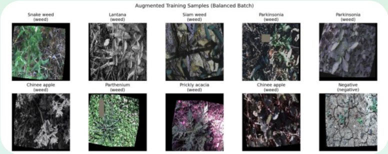

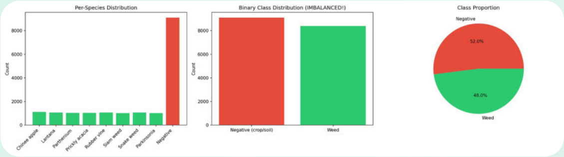

9 Classification Classes

8 weed species (~1,000 images each):

- Chinee apple
- Lantana
- Parthenium
- Prickly acacia
- Rubber vine
- Siam weed
- Snake weed
- Parkinsonia

Negative class: ~9,000 images (bare soil, grass, background)

Handling Class Imbalance

1

Class Weights in Loss Function

Formula: weight_i = total / (classes × count_i)

Minority classes receive higher weight, making errors on rare weeds cost more during training

2

Weighted Random Sampler

Each batch is balanced with minority classes oversampled for equal class representation

Ensures model sees all weed types frequently during training

CNN-ViT Hybrid Architecture

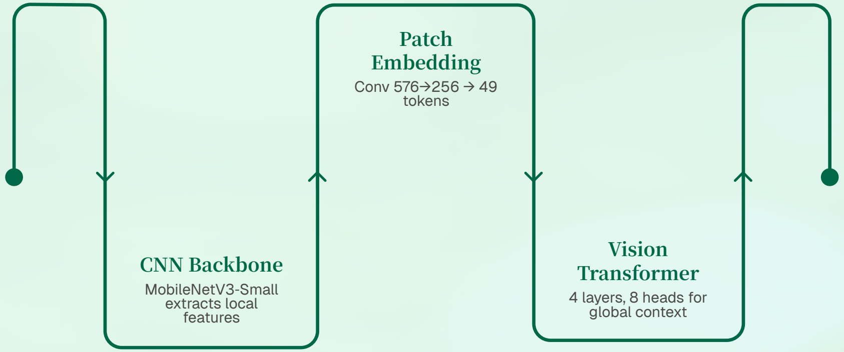

The hybrid architecture combines CNN's strength in local feature extraction with ViT's ability to capture global context, creating a lightweight model suitable for edge devices.

Why Hybrid Architecture?

- CNN excels at local features (texture, color, edges)
- ViT adds global understanding and context
- Lightweight design for edge devices
- Transfer learning: CNN pretrained on ImageNet, ViT trained fresh

Output Heads

Multi-class head: 9 species classification

Binary head: Weed vs. Negative decision

Combined outputs enable both species identification and spray decisions

Training Configuration & Results

Optimizer & Learning

AdamW optimizer with differential learning rates: CNN 3e-5 (fine-tune), ViT 3e-4 (train fresh)

Cosine scheduler, weight decay 0.01, gradient clip 1.0

Training Setup

Batch size 64, image resolution 224×224, 50 epochs with early stopping (patience 15)

Split: Train 70%, Val 15%, Test 15% (stratified)

Loss Function

Combined = 0.6 × MultiClass_CE + 0.4 × Binary_CE

Both losses use class weights to handle imbalance

Data Augmentation

Random crop, flip, rotation, color jitter, Gaussian blur, random erasing

---

Classification Performance

Test Set: 2,627 images

Multi-Class Results

81.73%

Accuracy

79.05%

F1 Macro

Best performers: Parthenium F1=0.927, Negative F1=0.869

Most challenging: Prickly acacia F1=0.612

Binary Classification

88.20%

Accuracy

88.44%

F1 Score

83.46%

Precision

94.05%

Recall

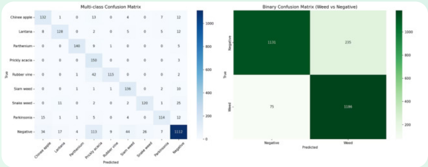

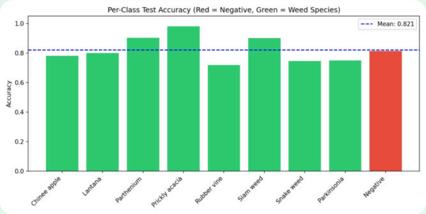

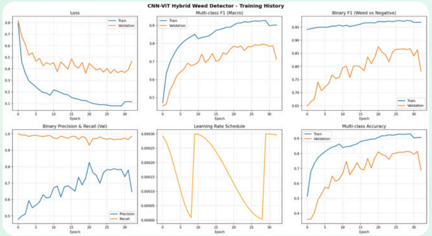

Key Insight: 94% recall catches almost every weed. 83% precision means some unnecessary spray but far better than 100% uniform spraying.

Ablation Study: CNN-ViT vs Pure CNN

Performance Comparison

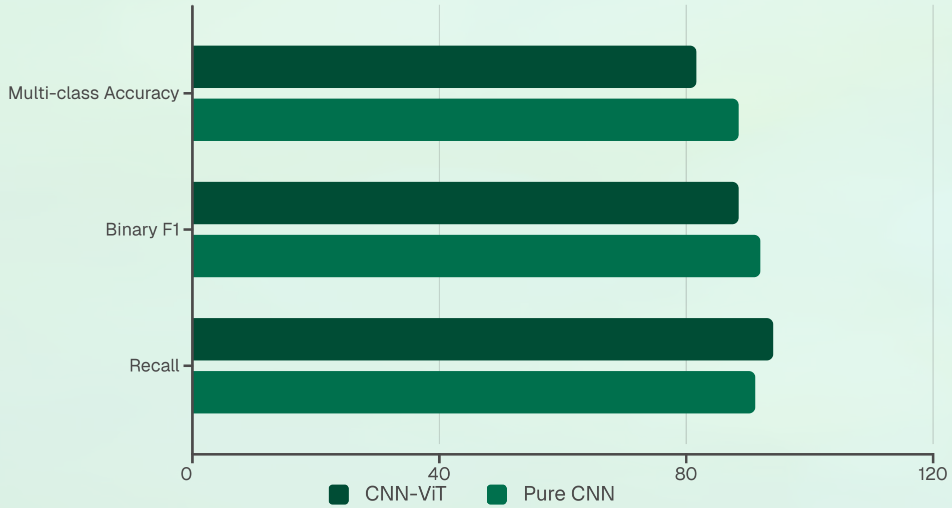

Finding: Pure CNN outperforms CNN-ViT by +3.5% Binary F1, but CNN-ViT achieves higher recall (94.05% vs 91.20%), catching more weeds with fewer misses.

Efficiency Metrics

Architecture	CNN-ViT	Pure CNN
Parameters	3.25M	1.91M
Model Size	12.4 MB	7.3 MB
Inference Time	6.5ms	5.1ms
FPS	306	389

Root Cause Analysis

1. **Dataset characteristics:** DeepWeeds contains single-plant close-ups (256×256). Classification depends on local features where CNN excels—ViT's global context is unnecessary
2. **Limited data:** 17.5K images relatively small. Transformers need more data; CNN's inductive biases work better with limited datasets

When ViT helps: Full-field aerial/drone images where spatial context matters, larger datasets (100K+ images), segmentation tasks

Scientific finding: For close-up plant images, CNN alone is sufficient. ViT shines on full-field images. This demonstrates understanding of WHY architectures work in different contexts.

Herbicide Reduction Impact

Simulation Results (2,627 tiles)

Traditional Approach

Spray all 2,627 tiles

100% coverage

AI-Guided Precision

Spray 1,421 tiles

54.1% coverage

45.9%

Herbicide Reduction

94.1%

Weed Catch Rate

1,186 of 1,261 weeds detected

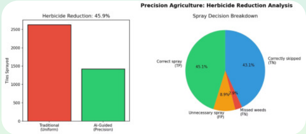

83.5%

Spray Efficiency

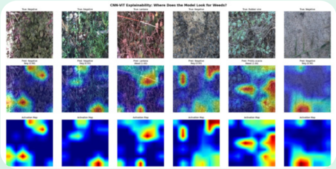

Spray Decision Thresholds

No Spray

Confidence < 0.3

Spot Spray

Confidence 0.3 - 0.7

Full Spray

Confidence > 0.7

Performance breakdown:

- 75 weeds missed (5.9%)
- 235 unnecessary sprays
- Significant reduction in environmental impact

Economic Impact

100-hectare farm example:

Traditional: ~~$10,000/year~~

AI-guided: **$5,410/year**

Annual savings: $4,590

Model Explainability with Grad-CAM

What is Grad-CAM?

Gradient-weighted Class Activation Mapping visualizes which regions of an image drive the model's decision. High gradient regions indicate areas that most influenced the prediction.

How it works: Grad-CAM uses gradients flowing back to the final convolutional layer to produce a coarse localization map highlighting important regions.

Red areas show where the model "looks" when making its classification decision.

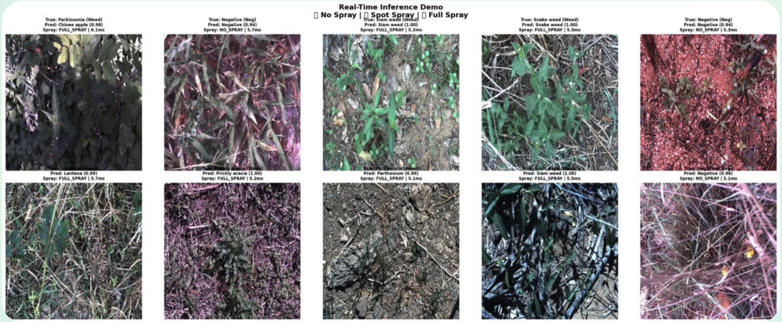

Key Findings

- ### Weed Images

  Model focuses on leaf and stem structures—the actual plant features that distinguish weed species

- ### Negative Images

  Diffuse activation with no focal point, confirming absence of distinctive plant features

- ### Validation

  Confirms model learned biologically relevant plant features, not background artifacts or spurious correlations

Importance for adoption: Explainability builds farmer trust. Grad-CAM demonstrates the AI makes decisions based on actual plant characteristics, not irrelevant background features. This transparency is crucial for real-world deployment.

Edge Deployment on Jetson Nano

Why Edge Computing?

- No internet connectivity required in field
- Real-time decision making
- Low power consumption
- Cost-effective ($99-149)

Deployment Pipeline

PyTorch (train, 12.4 MB) → ONNX (export, 12.5 MB) → TensorRT (deploy, ~6 MB FP16)

Performance Metrics

Kaggle CPU ONNX: 6.7ms / 149 FPS
Jetson Nano TensorRT: ~30-60ms / 15-30 FPS (estimated)

Optimization Technologies

ONNX Format

Open Neural Network Exchange

Universal format enabling train-anywhere, deploy-anywhere flexibility. Converts PyTorch models to platform-independent representation.

Maintains model accuracy while enabling cross-platform deployment.

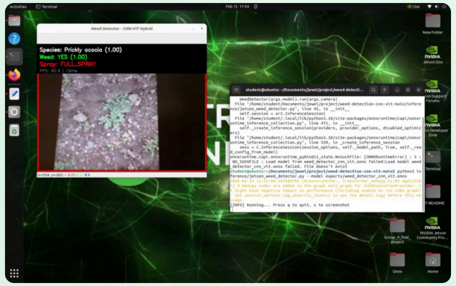

TensorRT Optimization

NVIDIA's inference optimizer

- Layer fusion for reduced memory access
- FP16 quantization (half precision)
- Memory optimization
- 2-3x speed improvement
- Model size: 12.4 MB → ~6 MB

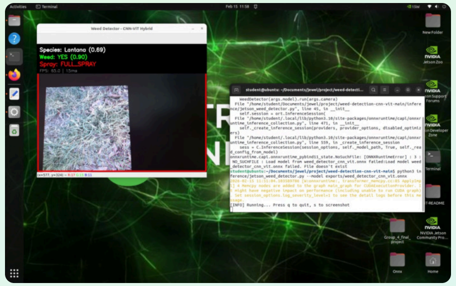

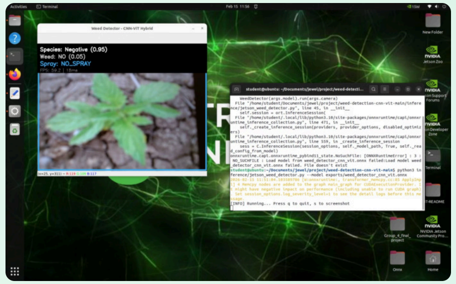

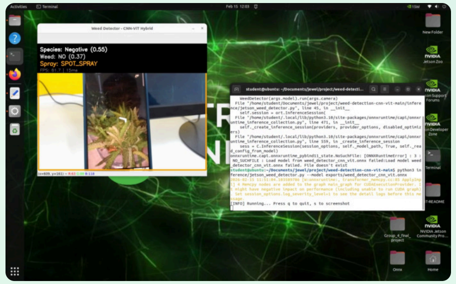

Real-Time Inference Pipeline

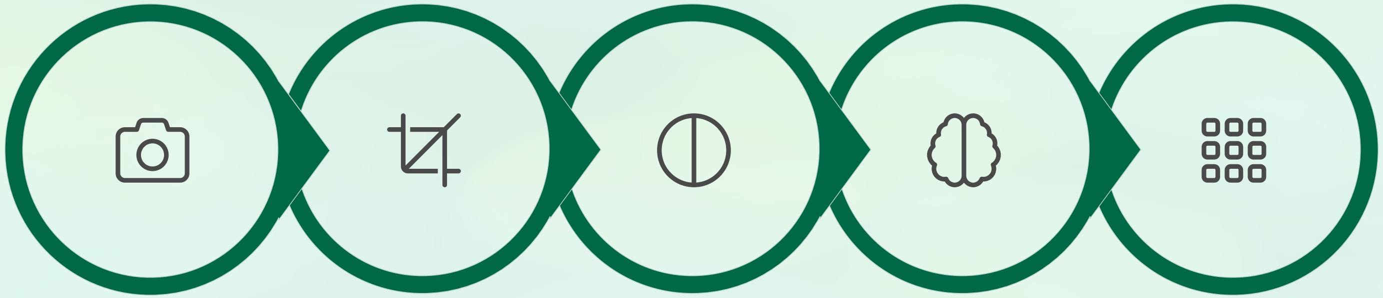

Frame capture

Resize 224×224

Normalize

Model inference

11 outputs

The complete pipeline processes camera frames in real-time, making spray decisions based on model confidence thresholds.

01

Setup

Flash JetPack, convert ONNX→TensorRT, connect USB camera

02

Execute

Run: python3 jetson_weed_detector.py

03

Deploy

Binary confidence >0.7 triggers spray decision

Limitations & Future Work

Current Limitations

Classification-Only Approach

No bounding boxes for individual plant localization. Tile-based approach is coarser than per-plant detection.

No Explicit Crop Class

Detects weed vs non-weed, not weed vs crop. Sufficient for spray decisions but not crop health monitoring.

ViT Shows No Improvement

Close-up images don't benefit from global attention mechanisms in this dataset.

Geographic Limitation

Single region (Queensland, Australia). May not generalize to other climates and weed species.

Future Research Directions

Object Detection

Use bounding box datasets (CWFiD, MFWD) for precise per-plant detection and localization

Field Validation

Deploy on actual Jetson Nano in agricultural settings for real-world performance validation

Crop Classification

Add crop-specific classes for broader applicability and crop health monitoring

Aerial Imagery

Test on full-field aerial images where ViT's global context advantage is pronounced

Hardware Integration

Integrate with sprayer nozzle controllers for complete end-to-end deployment system

Scientific maturity: Acknowledging limitations demonstrates research integrity. The architecture is generalizable—the same CNN-ViT hybrid can be retrained on any weed/crop dataset for any geographic region.